

Neural Network Viva Questions & Answers

1. Q: What is a Perceptron?

A: A perceptron is the simplest neural network model that takes inputs, applies weights, adds bias, and passes through an activation function to produce output. Example: Inputs [1,0], weights [0.5,0.3], bias=0.2 $\rightarrow z=0.7 \rightarrow \text{output}=1$ (step function).

2. Q: Why do we need a bias term in neural networks?

A: Bias allows shifting the activation function, improving flexibility of decision boundaries. Without bias, decision boundary passes through origin.

3. Q: What is the difference between Linear Regression and Perceptron?

A: Linear regression outputs continuous values, perceptron classifies with threshold. Example: Regression \rightarrow predict price, Perceptron \rightarrow classify spam or not.

4. Q: What is an activation function?

A: Introduces non-linearity in the model. Example: Sigmoid $\sigma(x)=1/(1+e^{-x})$. Input=2 \rightarrow output \approx 0.88.

5. Q: What is forward propagation?

A: Process of passing inputs through network to compute predictions.

6. Q: What is backpropagation?

A: Weights updated using gradient of loss w.r.t. weights via chain rule.

7. Q: What is the role of learning rate?

A: Controls step size in weight updates. High \rightarrow overshoot, Low \rightarrow slow training.

8. Q: What is a loss function?

A: Measures difference between predicted & actual values. Example: MSE = $\text{avg}((y_{\text{true}} - y_{\text{pred}})^2)$.

9. Q: Give an example of classification loss function.

A: Cross-Entropy Loss. Example: True=1, Pred=0.9 \rightarrow Loss= $-\log(0.9)=0.105$.

10. Q: What is an optimizer?

A: Algorithm used to minimize loss by updating weights. Example: SGD, Adam.

11. Q: What is bias in machine learning?

A: Error due to oversimplification. Example: Straight line fitting curved data.

12. Q: What is variance?

A: Error due to sensitivity to training data. Example: Overfitted tree good on train, poor on test.

13. Q: Explain bias-variance tradeoff.

A: More complexity → less bias but more variance.

14. Q: How to reduce overfitting?

A: Use regularization, dropout, early stopping, cross-validation.

15. Q: What is underfitting?

A: Model too simple to capture data. Example: Linear model on quadratic data.

16. Q: What is cross-validation?

A: Evaluating performance by splitting dataset into multiple train-test splits. Example: 5-fold CV.

17. Q: Why use cross-validation?

A: To avoid overfitting and estimate performance reliably.

18. Q: What is confusion matrix?

A: Table showing true vs predicted. Example: $\begin{bmatrix} 50 & 10 \\ 5 & 35 \end{bmatrix}$.

19. Q: Define Precision and Recall.

A: $\text{Precision} = \frac{TP}{TP+FP}$, $\text{Recall} = \frac{TP}{TP+FN}$. Example: $TP=80, FP=20, FN=10 \rightarrow P=0.8, R=0.89$.

20. Q: What is F1 score?

A: Harmonic mean of precision & recall. Example: $P=0.8, R=0.6 \rightarrow F1=0.685$.

21. Q: What is Gradient Descent?

A: Optimization method to minimize loss by moving opposite to gradient.

22. Q: Difference between SGD and Batch Gradient Descent?

A: SGD updates per sample, Batch after full dataset. Example: 1000 samples → 1000 updates in SGD, 1 update in Batch.

23. Q: Why is Adam optimizer popular?

A: Combines momentum & adaptive learning rate for faster convergence.

24. Q: What is Momentum?

A: Accelerates updates using past gradients. Like a rolling ball.

25. Q: What is RMSProp?

A: Uses moving average of squared gradients for adaptive learning rate.

26. Q: What is a feedforward neural network?

A: Info flows one-way (input→hidden→output).

27. Q: What is a CNN?

A: Specialized for image tasks using convolution & pooling layers.

28. Q: What is an RNN?

A: Has feedback loops, good for sequential data. Example: next-word prediction.

29. Q: What is LSTM?

A: RNN variant with memory cells to solve vanishing gradient problem.

30. Q: What is dropout?

A: Randomly drops neurons during training to prevent overfitting.

31. Q: What is vanishing gradient problem?

A: Gradients shrink to zero in deep nets, slowing learning.

32. Q: What is exploding gradient problem?

A: Gradients become too large, causing instability.

33. Q: What is batch normalization?

A: Normalizes inputs of each layer for stable training.

34. Q: Difference between supervised and unsupervised learning?

A: Supervised uses labeled data, unsupervised uses unlabeled. Example: classification vs clustering.

35. Q: What is reinforcement learning?

A: Learning with trial/error and rewards. Example: AI in chess.

36. Q: Train a perceptron on AND gate.

A: Inputs (0,0),(0,1),(1,0),(1,1). Outputs (0,0,0,1). Learned weights classify correctly.

37. Q: Explain bias & variance with example.

A: House price prediction: High bias → same prediction always. High variance → memorizes training noise.

38. Q: Why need non-linear activation?

A: Without it, network is just linear. Example: ReLU allows complex mapping.

39. Q: What is early stopping?

A: Stop when validation error rises though training error decreases.

40. Q: Why use cross-entropy loss in classification?

A: Measures distance between predicted vs true distribution. Example: True=[0,1], Pred=[0.2,0.8] → Loss=0.22.

41. Q: What is regularization?

A: Techniques to reduce overfitting by adding constraints.

42. Q: What is L1 regularization?

A: Adds absolute weights. Encourages sparsity.

43. Q: What is L2 regularization?

A: Adds squared weights. Prevents large weights.

44. Q: What is Elastic Net?

A: Combination of L1 and L2.

45. Q: How does dropout act as regularization?

A: Randomly drops neurons (output=0) during training.

46. Q: What is Early Stopping?

A: Stop when validation loss increases.

47. Q: What is Data Augmentation?

A: Expanding dataset with transformations.

48. Q: How does weight decay relate to L2?

A: Weight decay is another name for L2 regularization.

49. Q: Example where regularization is necessary?

A: Overfitting small dataset → apply L2 + dropout.

50. Q: How to choose λ in regularization?

A: Use cross-validation. Too high λ =underfit, too low=overfit.